

CS300 Couchbase NoSQL Server Administration

Lab 7 Exercise Manual



Release: 6.5.1

Revised: June 22nd, 2020



Lab #7: Advanced XDCR and Backup/Restore

Objective: This 1 hour lab will walk you through a couple of more advanced XDCR scenarios and then cover how to backup a bucket's files and restore it to a live cluster.

Overview: The following high-level steps are involved in this lab:

- Study optimistic replication in further detail
- Use pillowfight to insert 100,000 items UNDER the optimistic replication threshold into the 6-node NYC cluster so that the metadata is NOT checked on the NYC side before optimistically replicating it to the destination cluster
- Use pillowfight to insert 100,000 items OVER the optimistic replication threshold into the 6-node NYC cluster so that the metadata is checked on the NYC side before replicating it to the destination cluster
- Demonstrate that XDCR replication keeps occurring to remaining nodes even after a node goes down and fails over
- Learn how to use the vbucketkeygen tool to insert 1 key into each of the 1024 vBuckets
- Learn about cbbbackup and cbrestore
- Use cbbbackup to backup data in a bucket to some files
- Use cbrestore to restore the files back into a live bucket

Studying Optimistic Replication:

Remember that the Optimistic Replication threshold in our bidirectional replication stream is set to the default of 256 bytes. This parameter is basically used to tune the tradeoff between latency and bandwidth. If you want the fastest latency, you will have to sacrifice some bandwidth. On the other hand, if you want to conserve bandwidth, you will have to suffer some latency.

If the item that you write/update to either side cluster (NYC or London) is under 256 bytes, then the item is optimistically just thrown into the destination cluster's replication queue. However if the item you write/update is larger than 256 bytes, then the source cluster will first do a read of the same key's metadata over the WAN (Wide Area Network) to check if the destination cluster would win the conflict resolution. If the destination cluster's revision count is lower than the source cluster (where you did the write/update), then the source will send the item all the way down the WAN to the destination cluster. Before the destination cluster pushes that write/update to disk, it will do another metadata read locally to make sure that while the item was transiting over the WAN it didn't get updated with a higher revision count locally. If the destination cluster decides that the item it got from the source is still the higher revision count, then the destination cluster pushes that item into its local disk write queue.



Lab-7: Adv XDCR & Backup/Restore page 3

At this point, NYC-bucket and London-bucket should both be empty with zero items in either bucket.

Let's use pillowfight from the App Server to first send 100,000 keys of size 200 bytes (under the optimistic replication threshold) to the 6-node NYC cluster. In this case, the NYC cluster will NOT attempt to query metadata for any of the items from the remote side. It will just optimistically push all 100,000 items over to the London side and let London's cluster deal with it.

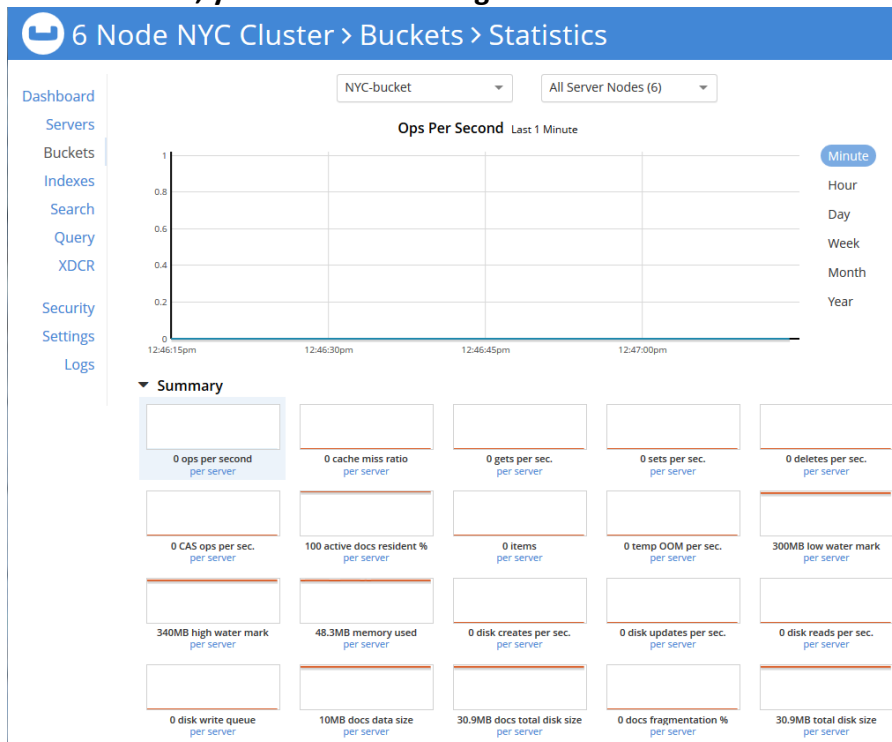
Later in the lab we will also send keys of size over 256 bytes to see how that has a different effect in the cluster.

Before we generate any data, let's set up our Web UI properly to see the effects of optimistic replication.

Switch browser tabs to the 6-node NYC cluster's Web UI, click on "Buckets link" at the side menu and click on the "statistics" to load the charts/graphs for this bucket(NYC-bucket):

6 Node NYC Cluster > Buckets							ADD BUCKET
Dashboard	name	items	resident	ops/sec	RAM used/quota	disk used	
Servers	gamesim-sample	586	100%	0	48.5MB / 400MB	31.6MB	Documents Statistics
Buckets	NYC-bucket	0	100%	0	48.3MB / 400MB	30.9MB	Documents Statistics
Indexes							

In the first tab, you should be seeing metrics for the "NYC-bucket" for "All Server Nodes":





Lab-7: Adv XDCR & Backup/Restore page 4

Next, switch to the 2nd browser tab for the London cluster's Web UI, click on "Buckets" at the side menu and click on the "statistics" to load the charts/graphs for the London-bucket:

Couchbase - London Cluster Enterprise Edition

Overview Server Nodes Data Buckets Query Indexes XDCR Security Log Settings

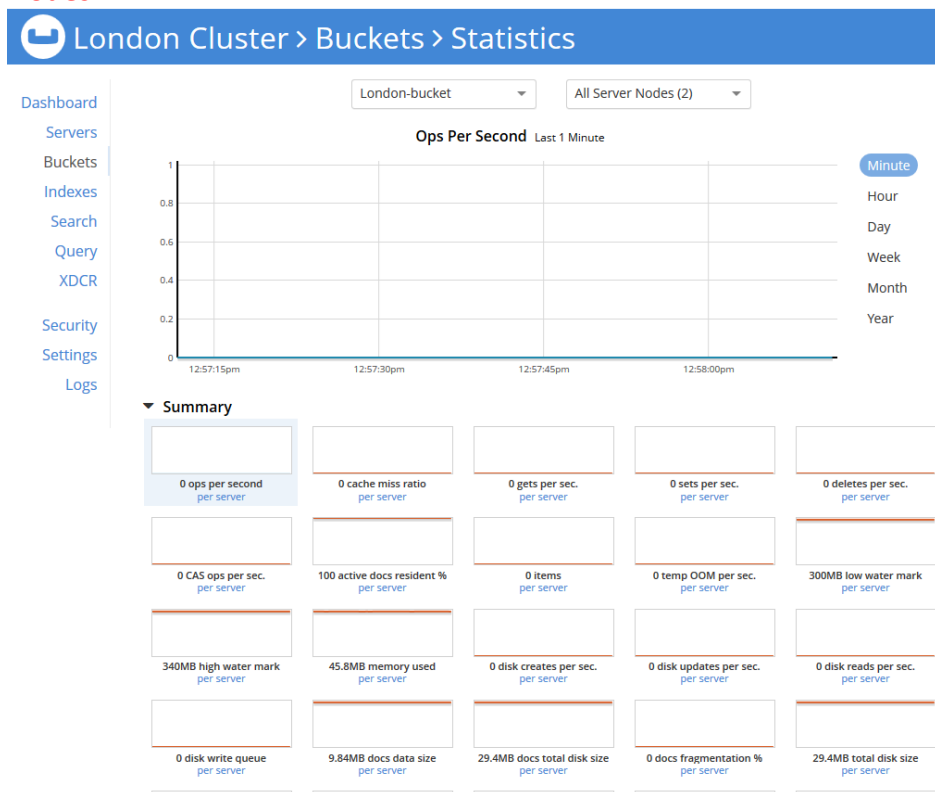
Data Buckets

Couchbase Buckets Create New Data Bucket

Bucket Name	Data Nodes	Item Count	Ops/sec	Disk Fetches/sec	RAM/Quota Usage	Data/Disk Usage
London-bucket	2	0	0	0	97.7MB / 400MB	9MB / 26.6MB


Documents Views

In the second tab, you should be seeing metrics for the "London-bucket" for "All Server Nodes":





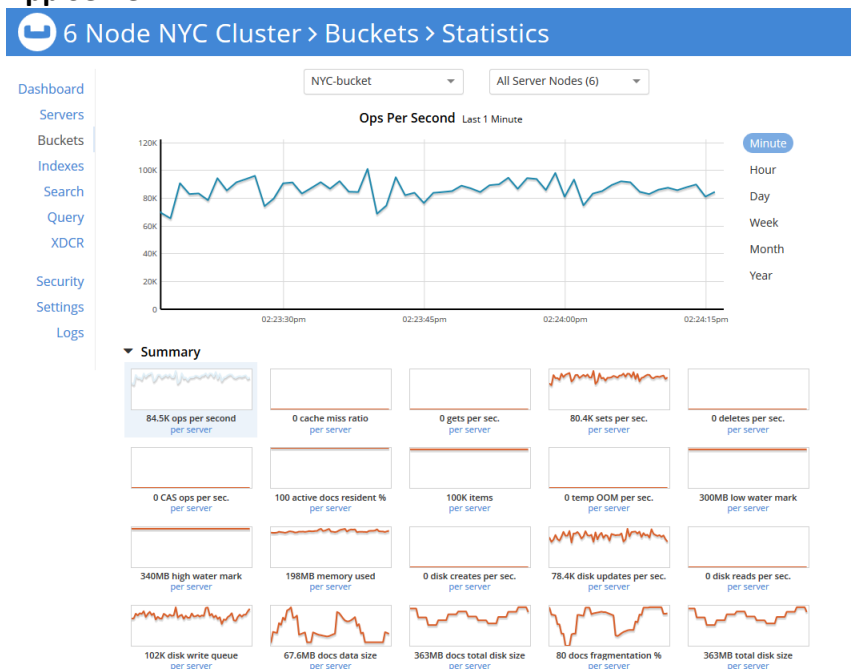
Switch to the App Server (black VM):



Run pillowfight to write 100,000 items, in 10 iteration, with a 100% write (set) ratio and a maximum payload size of 200 bytes. There will be a key prefix of “from_NYC” in front of all the keys. Use the public hostname of the 1st node in the 6-node NYC cluster in the command. So, we will push the writes into the 6-node NYC cluster and expect them to be optimistically replicated into the 2-node London cluster without first checking the metadata.

```
[ec2-user@appserver ~]$ cbc pillowfight -U couchbase://$NODE1/NYC-bucket --num-items=100000 --batch-size=10000 --set-pct=100 --min-size=50 --max-size=200 --timings --num-cycles=10 -p from_NYC -l
```

While the above command is generating data, **quickly switch to the 6-node NYC cluster’s Web UI** and observe that the cluster is receiving 40-70K writes per second across 4 nodes from the App Server:

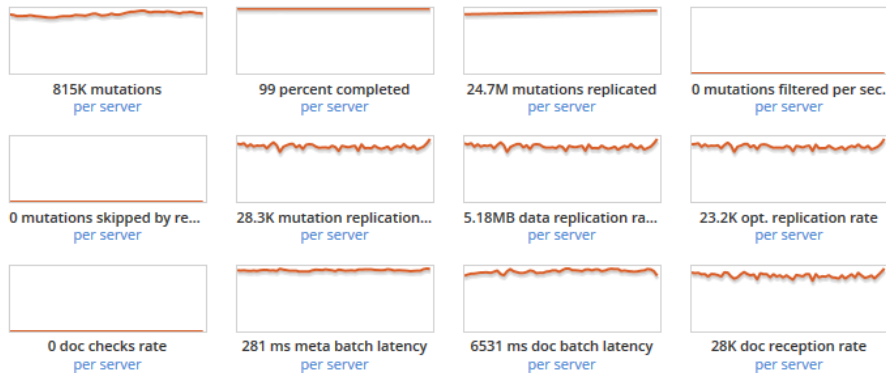


Then, in the same 6-node cluster’s Web UI, **scroll down and click the blue arrow to expand “OUTBOUND XDCR OPERATIONS”** to the London-bucket:



Lab-7: Adv XDCR & Backup/Restore page 6


▼ Outbound XDCR Operations to bucket "London-bucket" on remote cluster "London Cluster"



Observe in the screenshot above that the graph in the 2nd row, 4th column shows mutations replicated optimistically so far.

The graph in the 3rd row, 2nd column shows 250-300 millisecond latency for metadata operations.

Next, **switch to the 2nd browser tab for the 2-node London cluster** and in the large “ops per second graph” **observe that it is receiving around 30K operations (writes) per second from the NYC side:**





Scroll down on the same page, and expand “INCOMING XDCR OPERATIONS” and notice that the London cluster is getting around 25-30K sets per second (at least in my specific cluster at the time I captured the screenshot):


▼ Incoming XDCR Operations



The “metadata reads per sec” graph in the 1st column above is NOT referring to metadata reads occurring on the destination/London side, but rather metadata reads happening from source/NYC.

Next we’ll do a similar test of writing 100,000 keys into the NYC cluster with pillowfight, but this time we’ll set the item size to 500 bytes, well above the optimistic replication threshold. And actually we won’t write 100,000 new keys, but rather update the 100,000 keys from before (so we’ll use the same key prefix). This time the NYC cluster will have to keep doing remote metadata reads a total of 100,000 times before sending each item to the destination (London) cluster.

Switch to the App Server (black VM):



Run pillowfight to update the 100,000 items, in 10 iterations, with a 100% write (set) ratio and a maximum payload size of 10,000 bytes. There will be a similar key prefix of “from_NYC” in front of all the keys. Use the public hostname of the 1st node in the 6-node NYC cluster in the command. So, we will push the writes into the 6-node NYC cluster and expect 100,000 metadata reads to happen remotely from the 2-node London cluster before replicating the items over to the 2-node London cluster.

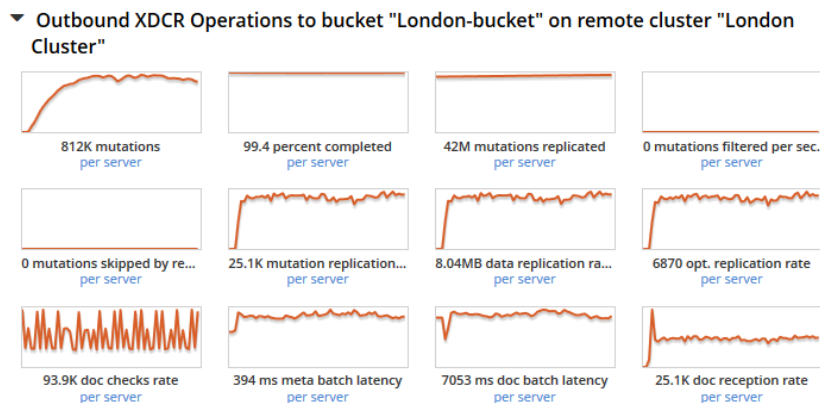


```
[ec2-user@appserver ~]$ cbc pillowfight -U couchbase://$NODE1/NYC-bucket --num-items=100000 --batch-size=10000 --set-pct=100 --min-size=50 --max-size=4000 --timings --num-cycles=10 -p from_NYC -l
```

Note that in the above command if you increase the item size to just barely over the optimistic threshold, like to 300 bytes, then optimistic replication won't technically kick in as this is an approximate setting. So, in our case, we are increasing the item size well over the 256 byte setting to 10,000 bytes to make sure that optimistic replication does not kick in at all for these keys. To be more technically accurate, for checking the optimistic replication threshold, Couchbase checks the compressed size of an item on disk when deciding whether to send it optimistically or not.

Remember, we are writing 100,000 keys with the exact same key prefix, so we are actually doing updates to the existing 100,000 keys from the previous run of pillowfight.

Anyway, switch to the 6-node NYC cluster's Web UI and look at the "OUTBOUND XDCR OPERATIONS" to the London-bucket:




Notice in the screenshot above that the graph in the 2nd row, 4th column shows lower optimistic rate of replication. This is because there are now more metadata reads occurring on the NYC side!

Switch to the 2-node London cluster's Web UI and look at the "INBOUND XDCR OPERATIONS":



Lab-7: Adv XDCR & Backup/Restore page 9



Observe here also that there are around 16-19K metadata reads per second happening. This is the 6-node NYC cluster doing metadata reads from London.

Stop the pillowfight command. (ctrl + C)

, remain in the 2-node London cluster's web UI and **click on "Buckets link"** at the side:

London Cluster > Buckets

name	items	resident	ops/sec	RAM used/quota	disk used
London-bucket	100,000	100%	0	141MB / 400MB	58.6MB

You will see a total of 100,000 items in the London-bucket. This is because we ran pillowfight in a loop with item count of 100,000 in this lab, the first 100,000 items written were actually updated.



Demonstrate that XDCR Replication keeps occurring to remaining nodes even after a node goes down and fails over:

Currently you should have bidirectional, encrypted replication set up between 6-node NYC and 2-node London. There are 100,000 keys with the prefix 'from_NYC' in buckets on both sides.

What do you think will happen if while you're inserting keys into the NYC cluster there is a node failure on one of the 2 nodes in London? If the 2nd node in London crashes, about half of the active vBuckets in London will be down and therefore half of the data from NYC can no longer be sent via XDCR to London.


But if we failover the node that crashed in London (2nd Node), then we would make the corresponding replicas on Node #1 in London active, thereby activating all 1024 vBuckets on the first London node.

In this section, you will then see that the items that could not be replicated temporarily while the 512 vBuckets in London were down, will suddenly start to be replicated to the newly promoted 512 vBuckets on the 1st node.

To demonstrate this concept we will use the vbucketkeygen tool to write 1 key to each of the 1024 vBuckets on the 6-node NYC cluster.

First, let's figure out how to use a combination of the Couchbase tools 'vbucketkeygen' and 'cbc create' along with the Linux tools curl, awk, echo and pipes to write exactly one key to each of the 1024 vBuckets in the NYC-bucket. This will essentially require us to write a simple script.

Switch to the App Server (black VM):



First pull the help menu for vbucketkeygen to learn how to use it:

```
[ec2-user@appserver ~]$ /opt/couchbase/bin/cbc-keygen --help
```

Usage:

```
keygen [OPTIONS ...]
```

Output a list of keys that equally distribute amongst every vbucket

```
-P --password          Bucket password [Default='']
-u --username          Username [Default='']
```



Lab-7: Adv XDCR & Backup/Restore page 11

```

-Z --config-cache          Path to cached configuration
[Default='']
-U --spec                  Connection string
[Default='couchbase://localhost/default']
--truststorepath           [Default='']
--certpath                 Path to server SSL certificate
[Default='']
--keypath                  Path to client SSL private key
[Default='']
-T --timings               Enable command timings
[Default=FALSE]
-v --verbose               Set debugging output (specify
multiple times for
r greater verbosity
[Default=FALSE]
--dump                     Dump verbose internal state after
operations are done [Default=FALSE]
-y --compress              Turn on compression of outgoing
data (second time to force compression)
[Default=FALSE]
-D --cparam <OPTION=VALUE> Additional options for
connection. Use -Dtimeout=t=SECONDS for KV operation
timeout [Default=]
--keys-per-vbucket         number of keys to generate per
vbucket [Default=1]
-? --help                  this message

```

Notice in the **yellow highlighted section** above, that you have to give this tool a mapfile, then 2 parameters: the keys per vbucket and how many keys to generate. This essentially means that we can generate 10,000 keys and then if the 'keys per vbucket' parameter is set to 1, we will randomly pick 1,024 of the 10,000 keys and push 1 key into each of the 1,024 vBuckets. This would NOT MEAN that we are going to distribute 10,000 keys evenly into 1,024 vBuckets! Instead, it would mean that we are going to distribute 1,024 random keys (from a seed pool of 10,000 keys) evenly into 1,024 vBuckets.

Run the following command to provide the output from curl as the mapfile and then generate 10,000 keys and place 1 of those keys into each of the 1024 vBuckets. Provide this command with the public hostname of Node #1 (dark blue VM in the 6-node NYC cluster).

```
SKIP THIS STEP>>>>>>>[ec2-user@appserver ~]$ curl -u Administrator:couchbase -n http://$NODE1:8091/pools/default/buckets/NYC-bucket | /opt/couchbase/bin/cbc-keygen 1
```



```
% Total      % Received % Xferd  Average Speed   Time    Time       Time
Current                                  Dload  Upload   Total   Spent    Left

Speed
100 12275  100 12275    0      0   328k      0 --:--:-- --:--:-- --:--:--
:-- 1712k
key_00000008638 0
key_00000001149 1
key_00000000146 2
key_00000001320 3
key_00000000093 4
key_00000001365 5
key_00000000418 6
<output truncated>
```

Notice in the above output that **key # ending 8638** would go into vBucket 0.

Now the above command does NOT actually push the keys into the vBucket. It is just a generator.

We have to pipe the output from the above command to series of other commands including 'cbc create' to actually push the keys into Couchbase.

In the following command, we are inserting 1 key into each of the 1,024 vBuckets in the bucket NYC-bucket in the 6-node NYC cluster. We are going to prefix each key with **AA.key#**. Prefixing with "AA." will lexographically sort all of the keys in the Web UI's view, so it'll be easy to find all the AA keys from this generation. Provide this command with the public hostname of Node #1 (dark blue VM in the 6-node NYC cluster).

Note: the command in red will work if using libcouchbase 2.8.2 or better

```
[ec2-user@appserver ~]$ for i in `curl -u Administrator:couchbase -n
http://$NODE1:8091/pools/default/buckets/NYC-bucket
|/opt/couchbase/bin/cbc-keygen -U couchbase://$NODE1/NYC-bucket --
keys-per-vbucket 1 | awk {'print $1'}`;do echo -n "value" |
/opt/couchbase/bin/cbc-create -U couchbase://$NODE1/NYC-bucket
AA.$i;done
```

```
AA.key_0000000877  Stored. CAS=0x158ead9bacd70000
AA.key_0000001878  Stored. CAS=0x158ead9bad980000
AA.key_0000001468  Stored. CAS=0x158ead9bae5b0000
AA.key_0000000139  Stored. CAS=0x158ead9baf220000
AA.key_0000009648  Stored. CAS=0x158ead9baf220000
AA.key_0000000350  Stored. CAS=0x158ead9bb0a30000
AA.key_0000000832  Stored. CAS=0x158ead9bb1610000
<output truncated>
```

If you now **switch to the browser tab for the 6-node NYC cluster and refresh the page**, you will see 101,024 keys in the NYC-bucket. We just added the last 1,024 keys.



6 Node NYC Cluster > Buckets			
Dashboard	name ▼	items	resident
Servers	gamesim-sample	586	100%
Buckets	NYC-bucket	101,024	100%
Indexes			

Try **clicking on Documents** for the “NYC-bucket” to actually display the items:

6 Node NYC Cluster > Buckets							ADD BUCKET
Dashboard	name ▼	items	resident	ops/sec	RAM used/quota	disk used	
Servers	gamesim-sample	586	100%	0	48.5MB / 400MB	31.6MB	Documents Statistics
Buckets	NYC-bucket	101,024	100%	0	125MB / 400MB	110MB	Documents Statistics
Indexes							

Here you will see the keys starting with “AA.” that we just added. **Change the dropdown for item count on the page in the bottom left corner from the default of 10 to 100:**

Dashboard

Servers

Buckets

XDCR

Security

Settings

Logs

Documents

Query

Search

Analytics

Eventing

Indexes

6 Node NYC Cluster > Documents

CLASSIC EDITOR ADD DOCUMENT

Bucket

NYC-bucket

Limit

10

Offset

0

Document ID

enter document ID or leave blank

Where

e.g., 'meta().id = "some_id" and 1

Retrieve Docs

10 Results for NYC-bucket, limit: 10, offset: 0

simple ☒ spreadsheet

< Prev Batch

Next Batch >

				AA.key_0000000000	Binary Document, dmFsdWU=
				AA.key_0000000001	Binary Document, dmFsdWU=
				AA.key_0000000002	Binary Document, dmFsdWU=
				AA.key_0000000003	Binary Document, dmFsdWU=
				AA.key_0000000004	Binary Document, dmFsdWU=
				AA.key_0000000005	Binary Document, dmFsdWU=
				AA.key_0000000006	Binary Document, dmFsdWU=
				AA.key_0000000007	Binary Document, dmFsdWU=
				AA.key_0000000008	Binary Document, dmFsdWU=
				AA.key_0000000009	Binary Document, dmFsdWU=

You may be surprised that all of the keys you see are in sequential order. For example the screenshot above shows keys 00 – 05 all in order. This does not mean that the script we ran inserts key 0 into vBucket 0 and key 24 into vBucket 24!



Try **clicking the page to page 10** and you will start to see that the AA. keys can get higher than 1024:

AA.key_0000009628	dmFsdWU=	Delete	Edit
AA.key_0000009629	dmFsdWU=	Delete	Edit
AA.key_0000009638	dmFsdWU=	Delete	Edit
AA.key_0000009639	dmFsdWU=	Delete	Edit


100 per page < page 10 >

Did the new 1,024 keys also get replicated to London? **Switch to the browser tab for the 2-node London cluster, refresh the page and look at the Data Buckets page to see if the item count has increased:**

London Cluster > Buckets				
	name	items	resident	ops/sec
	London-bucket	101,024	100%	0

It has increased. So replication is also working as expected.

Now we will do something interesting. Let's cause a simulated failure on the 2nd node in the London cluster. Log into the XDCR Node #2 in London (Orange VM/Couchbase08):



We're going to pretend that the Couchbase service has crashed here... or pretend that the entire machine has crashed. We'll simulate this by simply stopping the Couchbase service here:

```
[ec2-user@couchbase08 ~]$ sudo systemctl stop couchbase-server
Stopping couchbase-server
[ec2-user@couchbase08 ~]$ sudo systemctl status couchbase-server
couchbase-server is not running
```

Switch back to the Web UI for the 2-node London cluster and under the Data Buckets page, **notice that the bucket is now red**, signaling that half of the data is now unavailable b/c half of our London cluster is down.



Lab-7: Adv XDCR & Backup/Restore page 15

London Cluster > Buckets ADD BUCKET

name	items	resident	ops/sec	RAM used/quota	disk used	
London-bucket	50,512	100%	0	60.1MB / 400MB	24.5MB	Documents Statistics

1 node not responding

Type: Couchbase

Bucket RAM Quota: 400MB

Cluster RAM Quota: 2.07GB

Replicas: 1

Server Nodes: 2

Ejection Method: Value-Only

Conflict Resolution: Sequence Number

Compaction: Not active

Memory

cluster quota (4.14 GB)

other buckets (0 B)

this bucket (400 MB)

remaining (3.75 GB)

Disk

total cluster storage (19.9 GB)

other buckets (36.1 MB)

this bucket (24.5 MB)

remaining (16 GB)

Delete Compact Edit

Notice in the above screenshot also that the current Item Count is only 50,512 since half the active vBuckets are now missing! Keep this # in mind.

Click on Server link at the side and notice that the 2nd machine is in a down state:

London Cluster > Servers FILTER GROUPS ADD SERVER

Warning: Additional active servers required to provide the desired number of replicas.

name	group	services	CPU	RAM	swap	disk used	items	
ec2-54-89-225-199.compute-1.amazonaws.com	Group 1	data index query	8%	21.7%	---	24.5MB	50.5 K/50.5 K	Statistics
ec2-54-91-209-132.compute-1.amazonaws.com	Group 1	data index query	20.7%	15.3%	---	36.1MB	50.5 K/50.5 K	

Node unresponsive | Not taking traffic | FAILOVER to activate available replicas

Rebalance Failover

Switch back to the App Server (black VM). Now, let's try to write 1,024 keys into the 6-node NYC cluster and see how many of them successfully replicate to the 1-node London cluster. This time prefix the keys with "AB." b/c this is the 2nd time we're writing 1,024 keys. Still provide this command the public hostname of Node #1 in the 6-node NYC cluster.

```
[ec2-user@appserver ~]$ for i in `curl -u Administrator:couchbase -n
http://$NODE1:8091/pools/default/buckets/NYC-bucket
|/opt/couchbase/bin/cbc-keygen -U couchbase://$NODE1/NYC-bucket --
keys-per-vbucket 1 | awk {'print $1'}`;do echo -n "value" |
/opt/Couchbase/bin/cbc-create -U couchbase://$NODE1/NYC-bucket
AB.$i;done
```

```
AB.key_0000000285    Stored. CAS=0x158eae057590000
AB.key_0000000877    Stored. CAS=0x158eae057fc0000
AB.key_0000001878    Stored. CAS=0x158eae058980000
AB.key_0000001468    Stored. CAS=0x158eae059310000
AB.key_0000000139    Stored. CAS=0x158eae059cd0000
AB.key_0000009648    Stored. CAS=0x158eae05a670000
```



Lab-7: Adv XDCR & Backup/Restore page 16

```
AB.key_0000000350    Stored. CAS=0x158eae05b020000
AB.key_0000000832    Stored. CAS=0x158eae05b9d0000
<output truncated>
```

Switch to the browser tab for the 6-node NYC cluster, go to the Data Bucket page, refresh it and notice that the item count is 102,048. This makes sense. Originally the bucket had 100,000 keys (from the pillowfight command we ran in the prior section). Then we added 1,024 keys starting with AA. The new added 1024 more keys starting with AB.

6 Node NYC Cluster > Buckets			
	name ▼	items	resident
	gamesim-sample	586	100%
	NYC-bucket	102,048	100%

Switch to the browser tab for the 2-node London cluster (which is down to 1 healthy node), go to the Bucket page, refresh it and notice that the item count is 51,024.

This makes sense also. Remember that after we lost 1 node in London, the item count went down from 101,024 to 50,512 (it became half). Then we replicated 512 or so more keys with the prefix AB from NYC -> London. So $50,512 + 512 = 51,024$. The issue here is that there are 512 keys in a sort of “pending replication” state in the NYC cluster. The NYC cluster cannot send them to London b/c half the active vBuckets in London are missing. We have to Fail Over the down node in London to promote the replicas on the remaining node. Then the updated cluster map from London (with one node owning all 1,024 vBuckets) will be sent to all nodes in NYC and the 4 NYC nodes will push out the last 512 keys with prefix AB over to the London 1-node cluster.

Note that a Rebalance is not required in London to get the cluster map updated on the NYC side.

London Cluster > Buckets							ADD BUCKET
	name ▼	items	resident	ops/sec	RAM used/quota	disk used	
	London-bucket	51,024	100%	0	60.1MB / 400MB	28.5MB	Documents Statistics

1 node not responding

Switch back to the London cluster’s Web UI, click on “Servers” at the side, then click the “Fail Over” button next to the down node to promote its replicas on the remaining node.



Remember that we configured the London-bucket with 1 replica (so 2 copies, one active and one replica).

Confirm Failover - ec2-54-91-209-132.compute-1.amazonaws.com:8091 X

Attention: There are not replica (backup) copies of all data on this node! Failing over the node now will irreversibly lose that data when the incomplete replica is activated and this node is removed from the cluster. If the node might come back online, it is recommended to wait. Check this box if you want to failover the node, despite the resulting data loss.

☒ Confirm failover

[Cancel](#) [Failover Server](#)

Side note: The Fail Over cannot be configured to occur automatically in London, since there are only 2 nodes in the cluster. There have to be at least 3 nodes in a Couchbase cluster before auto-failover can be enabled. This is because with a 2-node cluster, if failover was automatic, then a split-brain condition could have occurred. (Well, technically it is possible to configure auto-failover in a 2-node cluster, but it won't get triggered if a node goes down. If you grow the same cluster to 3+ nodes, then auto-failover will start to get triggered when a node goes down.)

On the pop-up, place a check next to “Please confirm Failover” and click “Fail Over”:

Now the GUI is saying that a Removal will happen at next rebalance.

London Cluster > Servers FILTER GROUPS ADD SERVER									
Warning: At least two servers with the data service are required to provide replication.									
Dashboard	name	group	services	CPU	RAM	swap	disk used	items	Rebalance
Servers	ec2-54-89-225-199.compute-1.amazonaws.com	Group 1	data index query	5.05%	21.6%	---	30.5MB	101 K/0	Statistics
Buckets	ec2-54-91-209-132.compute-1.amazonaws.com	Group 1	data index query	20.7%	15.3%	---	36.1MB	50.5 K/50.5 K	
Indexes	Node unresponsive, failed-over Not taking traffic REMOVAL pending rebalance								
Search									
Query									

At this point, the updated cluster map from London (that one node owns all 1,024 vBuckets) has been sent to all 4 data nodes in NYC and the 4 data NYC nodes have started pushing out the last 512 keys with prefix AB over to the London 1-node cluster.

In this case, the Web UI is simply asking you to run a Rebalance to remove down node from the cluster. If you had clicked on Rebalance, no rebalancing would actually occur. The rebalance is being requested to permanently remove the down node from the vBucket cluster map.

Let's verify this.

In the London (now 1-node) cluster's Web UI, click on Buckets link at the side and look at the item count:



Lab-7: Adv XDCR & Backup/Restore page 18

The item count now matches what we see in the 6-node NYC cluster. Perfect! *(If you don't see the same item count yet, you may need to wait 1 – 2 mins for XDCR to finish replicating and then refresh the page)*

This proves that as soon as you Failed Over the down node in London, the 6-node NYC cluster has the 512 vBuckets become active again on the remaining London node and then the 4 nodes in NYC started to replicate out those last 512 keys to London.


The steps we completed above demonstrate that if we had auto-failover configured in London with 3+ nodes in the London cluster, then even if we lose a node, replication would eventually resume automatically after the auto-failover is triggered.

Before concluding this section, let's Rebalance the London cluster so that the Web UI shows only 1 node, then we'll re-add the down node back into the cluster and do another rebalance so that the London cluster becomes a healthy 2 node cluster again.

Switch to the browser tab for the London cluster (currently 1 node), and click on "Server Nodes" at the top and then hit Rebalance:

The rebalance will finish in just a few seconds as there is really nothing to rebalance:

Log into the XDCR Node #2 in London (Orange VM):



This is the node in London that we had previously stopped the Couchbase service on. **Verify that the Couchbase service is still stopped and then start it again:**

```
[ec2-user@ip-172-31-41-211 ~]$ sudo systemctl status couchbase-server
couchbase-server is not running
```

```
[ec2-user@ip-172-31-41-211 ~]$ sudo systemctl start couchbase-server
Starting couchbase-server [ OK ]
```

Add the server back to the two node London cluster using the Join method covered in Lab #3

Couchbase > Join Cluster

Cluster Host Name/IP Address

Cluster Admin Username

Cluster Admin Password

▼ Configure Services & Settings For This Node

☒ Data Service

☒ Index Service

☐ Search Service

☒ Query Service

This Node: Host Name/IP Address Usually localhost or similar

Data Disk Path Path cannot be changed after setup

Free: 8 GB

Indexes Disk Path Path cannot be changed after setup

Free: 8 GB

[< Back](#) [Join With Custom Configuration](#)



Lab-7: Adv XDCR & Backup/Restore page 20

You will see a #1 under Pending Rebalance now as the new node has been added to the cluster, but the active and replica vBuckets still have to be redistributed across the two nodes. **Click on Rebalance:**

London Cluster > Servers

Warning: At least two servers with the data service are required to provide replication.

✓ This server has been associated with the cluster and will join on the next rebalance operation.

Rebalance

name	group	services	CPU	RAM	swap	disk used	items
ec2-89-225-199.compute-1.amazonaws.com	Group 1	data index query	6.06%	23.2%	---	36.5MB	102 K/0
ec2-54-91-209-132.compute-1.amazonaws.com	Group 1	data index query	1.5%	13.1%	---	---	0/0

New node | Not taking traffic | ADD pending rebalance

Cancel Add

The rebalance will start and should take about 2 – 3 minutes to complete:

London Cluster > Servers

rebalancing 2 nodes 16.4%

Stop Rebalance

name	group	services	CPU	RAM	swap	disk used	items
ec2-54-89-225-199.compute-1.amazonaws.com	Group 1	data index query	78.9%	23.5%	---	36.5MB	90.9 K/11 K
ec2-54-91-209-132.compute-1.amazonaws.com	Group 1	data index query	27.6%	14.6%	---	2.79MB	12.4 K/101

8.37% complete

24.5% complete

This concludes Lab #7.